

## CS552: Generative AI, Homework 4

Will Buchta

3/25/2025

1. **Diffusions: Unconditional Generation [13 pts]:** Using a trained diffusion model (you can either just use the pre-trained DDPM in `diff_unet_faces.cpt` or, if you prefer, train your own using the face data in `faces23k.48x48.npy.gz`), sample 100 faces (each is 48x48 pixels) and display them in a 10x10 collage. For DDPM, the sampling procedure is given in the Prince textbook (linked on Canvas), Algorithm 18.2. For  $\sigma_t$ , a common practice (according to the original DDPM paper by Ho et al, 2020) is to set  $\sigma_t = \sqrt{\beta_t}$ . The hyperparameters used for training, including the noise schedule  $\{\beta_t\}_{t=1}^T$ , are given in the starter code `homework4.template.py`.

See Figure 1 for my results and `ddpm.py` for the implementation.



*Figure 1: Unconditional face diffusion*

2. **Diffusions: Inpainting/Merging [12 pts]:** Using the same pre-trained DDPM from the previous problem, “merge” two faces into one:

- Select any two faces  $\mathbf{x}^{(1)}$  and  $\mathbf{x}^{(2)}$ .
- Apply the diffusion kernel for some  $t < T$  to compute  $\mathbf{z}_t^{(1)}$  and  $\mathbf{z}_t^{(2)}$ .
- Create  $\mathbf{z}_t$  by combining the left half of  $\mathbf{z}_t^{(1)}$  with the right half of  $\mathbf{z}_t^{(2)}$ .
- Sample  $\mathbf{x} \sim P(\mathbf{x} \mid \mathbf{z}_t)$ .

Apply this to 5 pairs of faces each for 4 different values of  $t$  each – show the results in a 4x15 collage (where the rows correspond to different  $t$  – 10 columns show the 5 pairs of  $(\mathbf{x}^{(1)}, \mathbf{x}^{(2)})$ , and 5 columns show the corresponding merged samples  $\mathbf{x}$ . Arrange the columns as you see fit but make sure to explain them clearly in your report.

See problem2() in ddpm.py. For each group of three columns,  $\mathbf{x}^{(1)}$  is the first photo,  $\mathbf{x}^{(2)}$  is the second photo, and  $\mathbf{x}$  is the third photo. The best results can be seen when the diffusion kernel is applied 500 times in Figure 2:



Figure 2: Merged faces

3. **Diffusions: Conditional Generation with Classifier Guidance [30 pts]:** Using the face data in `faces23k_48x48.npy` and associated age labels in `ages23k.npy`, train an age regressor  $f$  that maps from any  $(\mathbf{z}_t, t)$  to an estimate  $\hat{y}$  of how old the person in the corresponding  $\mathbf{x}$  is. To create the training data for  $f$ , you will need to use the diffusion kernel to create  $\mathbf{z}_t$  for various  $t$  and  $\mathbf{x}$ . Exclude any examples from training if either  $y < 0$  (which indicates a missing label) or  $y > 100$  (which is likely an incorrect label).  $f$  should output the mean of a Gaussian distribution with unit variance that characterizes the age of a specific noised image at a specific timestep  $t$ , i.e.,  $P(y | \mathbf{z}_t) = \mathcal{N}(f(\mathbf{z}_t, t), 1)$ . As an architecture for the regressor, you can use the first half (i.e., downward convolution) of the same U-Net with time embeddings which implements the noise estimator  $E_\theta$  in the pre-trained DDPM. However, the choice is up to you.

After training the regressor, use classifier guidance (even though you actually have a regressor) to generate 4 faces each for target ages 18, 40, 60, and 80 years old. To compute  $\frac{\partial \log P(y | \mathbf{z}_t)}{\partial \mathbf{z}_t} = -\frac{\partial (\hat{y} - y)^2}{\partial \mathbf{z}_t}$  for a target age value  $y$ , you can ask ChatGPT for help with the autograd expressions (you will need to enable `requires_grad=True` on the tensor representing  $\mathbf{z}_t$ ). You may find it useful to multiply the gradient by a small number (e.g., 0.2) rather than using its full magnitude. Show these faces in a 4x10 collage.

See my implementation in `ddpm.py`. I had to use a very small ‘guidance scale’ for the classifier’s gradient – see below for the difference between a scale of 0.001 and 0.0005. Any values higher than that resulted in “mush” for lack of a better word. The rows of the figures correspond to ages 18, 40, 60, and 80, which I think was represented quite well.



Figure 3: Classifier guidance with a scale of 0.001



Figure 4: Classifier guidance with a scale of 0.0005